

Writing Testable Code

F17

Principles

- ▶ Logically (albeit simplistic), we can think of a test as $y = f(x)$
- ▶ A test case is thus a tuple (x, e) where x is the input and e the expected result
- ▶ This works well for the “leaf components”/bottom layers of the system, but gets more complicated when we move up the layer stack
- ▶ For example, what about $y = g(u)$ where $u = f(x)$?

(We could imagine that g above is a logger and f produces a time stamp.)

Principles (cont'd)

Again, we are testing a layered/composite function $y = g(u)$ where $u = f(x)$

- ▶ If the test cases are composed into the form (x, e) , then where is the error?
- ▶ Locating an error (blame control) is extremely important
- ▶ Composition = integration testing, which is good but on a higher level
- ▶ So, how can we test g in isolation?
- ▶ We want (u, e) and not care about f

Principles (cont'd)

Consider the test program t , we can do either

- ▶ $t(x) = g(u)$ where $u = f(x)$, and
- ▶ $t(f, x) = g(f(x))$

What's the difference?

Principles (cont'd)

- ▶ $t(x) = g(u)$ where $u = f(x)$
 - ▶ The dependency to f is fixed inside the implementation of t . We can't change it or test g in isolation.
- ▶ $t(f, x) = g(f(x))$
 - ▶ f is passed as a dependency
 - ▶ Easy to test g in isolation (take f out of the equation by using test cases where the relation between $u = f(x)$ are known)
- ▶ This is higher-order programming (n.b. an object is a higher-order function)
- ▶ Extracting the behaviour breaks long chains of dependence that make components hard to reuse
- ▶ Factored out dependencies gives smaller building blocks; again easy to reuse

Good code is testable code

When writing code, we must *always* consider its *testability*

- ▶ Different functions should be testable in isolation
 - ▶ Minimise dependencies
 - ▶ Minimise possible sources of error
- ▶ Encapsulation can hamper testability
 - ▶ Can all functions be tested/accessed?
 - ▶ Can we easily supply values to test?
 - ▶ Can be easily get control input/output?
 - ▶ Code coverage vs. encapsulation?

Example 1: Logger

Let's say we have designed and implemented a library for logging C-strings to disk in `logger.c`

- ▶ The logger must be initialised in `initLogger(filename)`
- ▶ Messages are written to the logger with `logMessage(msg)`
- ▶ The logger is torn down with `destroyLogger()`
- ▶ Messages are buffered and flushed internal to the logger

Let's have a look at the code, and then talk about its testability.

```
/* logger.c */

#include <stdio.h>
#include <assert.h>
#include <string.h>
#include <time.h>

#include "logger.h"

/* Constants */
#define BUFSIZE 1048576
#define TIMESTAMPMAX 26

/* Module variables */
static FILE *logfile;
static char logbuffer[BUFSIZE];
static unsigned int logsiz = 0;
static time_t logtime;

/* Start code */
void initLogger(const char *fn) {
    assert(fn);
    assert(logfile == NULL);

    logfile = fopen(fn, "w");
}
```



```

void logMessage(const char *msg) {
    assert(msg);
    assert(logfile);

    int msgsiz = strlen(msg) + 1;

    if (logsiz + msgsiz + TIMESTAMPMAX > BUFSIZE) {
        flush();
    }

    time(&logtime);
    char *timestamp = ctime(&logtime);

    while (*timestamp != '\n')
        logbuffer[logsiz++] = *timestamp++;
    logbuffer[logsiz++] = ' ';
    while (*msg)
        logbuffer[logsiz++] = *msg++;
    logbuffer[logsiz++] = '\n';
    logbuffer[logsiz] = '\0';
}

static inline void flush() {
    fwrite(logbuffer, logsiz, 1, logfile);
    logsiz = 0;
}

```

```
void destroyLogger() {
    assert(logfile);

    flush();
    fclose(logfile);

    logfile = NULL;
}

/* Sample use
int main(void) {
    initLogger("testlog.txt");
    logMessage("Foo");
    logMessage("Bar");
    destroyLogger();
    return 0;
}
*/
```

Can you see any problems with testing the logger?

Let's examine what tests we should write

- ▶ Are log messages printed correctly?
- ▶ Are log messages' time stamps correct?

Problem list

1. The logger always writes to a file on the disk
2. The size of the buffer is fixed and requires a somewhat large test data
3. Time stamps are created internally so comparing two tests not possible (will always differ on time)

This design is not so good – the code is difficult to test, reuse and modify

An improved logger

1. Control where the output goes
2. Control over buffer size
3. Control of time stamps

```

/* better-logger.c */

#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <string.h>
#include <time.h>

/* Constants and helper functions */
#define TIMESTAMPMAX 26
#define initLogger(p) initLoggerWithPath(p, BUFSIZE)

/* Module variables */
static FILE *logstream;
static char *logbuffer;
static unsigned int logsiz;
static time_t logtime;
static unsigned int BUFSIZE = 1048576;
static unsigned int USES_BUFFER = 1;

void initLoggerWithPath(const char *_fn, unsigned int _bufsiz) {
    assert(_fn);
    assert(logstream == NULL);

    initLoggerWithStream(fopen(_fn, "w"), _bufsiz);
}

```

```

void initLoggerWithStream(FILE *_logstream, unsigned int _bufsiz) {
    assert(_logstream);
    assert(logstream == NULL);

    BUFSIZE = _bufsiz;
    logstream = _logstream;

    if (USES_BUFFER = BUFSIZ > 0) {
        logbuffer = (char*) malloc(BUFSIZE);
    }
}

static inline void flush() {
    fwrite(logbuffer, logsiz, 1, logstream);
}

static inline void flushAndReset() {
    flush();
    logsiz = 0;
}

const char *buffer() {
    return logbuffer;
}

```



```

void logMessageWithTime(const char *_msg, const time_t *_logtime) {
    assert(_msg);
    assert(_logtime);

    const int msgSize = strlen(_msg) + TIMESTAMPMAX;

    if (USES_BUFFER) {
        if (logsiz + msgSize > BUFSIZ) flushAndReset();
    } else {
        if (BUFSIZ < msgSize) logbuffer = realloc(logbuffer, BUFSIZ = msgSize);
    }

    char *timestamp = ctime(_logtime);

    while (*timestamp != '\n')
        logbuffer[logsiz++] = *timestamp++;
    logbuffer[logsiz++] = ' ';
    while (*_msg)
        logbuffer[logsiz++] = *_msg++;
    logbuffer[logsiz++] = '\n';
    logbuffer[logsiz] = '\0';

    if (!USES_BUFFER) flushAndReset();
}

```

```

void logMessage(const char *_msg) {
    assert(_msg);

    time(&logtime);
    logMessageWithTime(_msg, &logtime);
}

void destroyLogger() {
    assert(logstream);

    flushAndReset();
    fclose(logstream);
    free(logbuffer);

    logstream = NULL;
}

/* Sample use
int main(void) {
    initLogger("/dev/null");
    logMessage("Foo");
    fprintf(stderr, "In buffer: %s", buffer());
    destroyLogger();
    return 0;
}
*/

```

Observations

- ▶ Testing the logging isolated from file handling is now possible
- ▶ Code slightly more complex, but mostly just does the “obvious things” and depends on library functions being correct
- ▶ The code is now easy to test, reuse and extend
- ▶ (But what about timestamps?)

```

from time import ctime

def buffered(time, msg):
    global logbuffer
    if len(logbuffer) > 1024:
        for msg in logbuffer:
            print msg
        logbuffer = [time + " " + msg]
    else:
        logbuffer.append(time + " " + msg)

def unbuffered(time, msg):
    print time, msg

def initLogger(behaviour = buffered):
    global logbehaviour, logbuffer
    logbehaviour = behaviour
    if behaviour is buffered:
        logbuffer = []

def logMessage(msg, time = ctime()):
    logbehaviour(time, msg)

initLogger()
logMessage("Foo")
logMessage("Bar")

```

But even better

We can create a behaviour that simply prints to a specified string and that's that.

```
def createStoreToArrayBehaviour(buffer):  
    def storeToBuffer(time, msg):  
        buffer.append([time, msg])  
    return storeToBuffer  
  
myBuffer = []  
initLogger(createStoreToArrayBehaviour(myBuffer))  
logMessage("Foo")  
logMessage("Bar")  
for msg in myBuffer:  
    for element in msg:  
        print element
```

Which prints the elements as expected.

Testing a binary search tree

```
/* bst.h */

typedef struct _tree tree;
typedef tree *Tree;
struct {
    int value;
    Tree left, right;
};

Tree mkTree(int v);
void insert(Tree t, int v);
```

Testing a binary search tree (cont'd)

What's wrong with this test and/or module?

```
Tree t = mkTree(5);  
insert(&t, 1);  
insert(&t, 3);  
insert(&t, 7);  
assert(t->element == 5)  
assert(t->left->element == 1)  
assert(t->left->left == NULL);  
...
```

Testing a binary search tree (cont'd)

```
/* Returns the number of nodes in a tree */
int size(Tree t) {
    return (t) ? 1 + size(t->left) + size(t->right) : 0;
}

/* Returns the longest path to a leaf in a tree */
int depth(Tree t) {
    return (t) ? 1 + max(depth(t->left), depth(t->right)) : 0;
}
```

Allows us to test important properties of a tree without knowing about its implementation. Does it grow on insert? Of duplicates?

```
Tree t = mkTree(1);
assert(depth(t) == size(t) == 1)
for (int i=2; i<5; ++i) {
    insert(&t, i);
    assert(depth(t) == size(t) == i)
}
```


Testing a binary search tree (cont'd)

```
char *getPathForElement(Tree t, int element) {
    char *result = *path = (char*) malloc(depth(t));
    while (t) {
        if (element == t->element) {
            *path = '\0';
            return result;
        } else if (element < t->element) {
            *path++ = 'L';
            t = t->left;
        } else {
            *path++ = 'R';
            t = t->right;
        }
    }
    return NULL;
}
```

Allows us to trace e.g., moving elements on delete, etc.

```
Tree t = mkTree(3); insert(&t, 5); insert(&t, 1); insert(&t, 2);
assert(strcmp(getPathForElement(t, 2), "LR") == 0)
assert(strcmp(getPathForElement(t, 5), "R") == 0)
```

Testing a binary search tree (cont'd)

```
int getElementForPath(Tree t, char *path) {
    while (t && *path) {
        switch (*path++) {
            case 'L':
                t = t->left;
                continue;
            case 'R':
                t = t->right;
                continue;
            default:
                return -2;
        }
    }
    return t ? t->element : -1;
}
```

Allows us to trace e.g., moving elements on delete, etc.

```
Tree t = mkTree(3); insert(&t, 5); insert(&t, 1); insert(&t, 2);
assert(getElementForPath(t, "LR") == 2)
assert(getElementForPath(t, "R") == 5)
```

Observations

- ▶ Writing testable code will force you to stay away from certain patterns
 - ▶ Example: function that initialises an entire data structure in a single hit (constructor)
 - ▶ We want to be able to do piecemeal testing
 - ▶ Downside: can observe object in invalid state
- ▶ Avoid global state
 - ▶ Persists between tests (includes singletons)
 - ▶ No global state sometimes makes things more complex
- ▶ **Do use** the single responsibility pattern
 - ▶ Consequence: more units of code (functions, classes, modules, etc.)
- ▶ Minimise dependencies
 - ▶ ... or your tests will easily become very complex

Hints

- ▶ Picking good names is extremely important
- ▶ Be smart but avoid “fancy coding”
- ▶ Short code is readable code, too short code is unreadable
- ▶ Many small functions that can be composed increases reusability and maintainability – and therefore testability!
- ▶ (Think of testing as reuse)
- ▶ Use asserts, especially for “this should never happen”
- ▶ Avoid `NULL`
- ▶ Always initialise variables, even if you “know” they will be assigned before read later
- ▶ Remove redundant or unused things – less clutter is more readable
- ▶ Avoid assignments in boolean expressions and function arguments

Hints (cont'd)

- ▶ Test for division by zero if you don't know the value of the denominator
- ▶ Make the loop invariant clear – preferably in a single place
- ▶ Avoid many **breaks**, **returns** etc. – makes code hard to follow
- ▶ Always test for the most likely case first (that's the one the next guy will be looking for anyway)
- ▶ Always test for boundary conditions and array sizes
- ▶ Inspect the return values for functions you expect to always succeed (like `malloc`)
- ▶ Whenever you use `malloc` inside a function to return data, note it in the function's documentation
- ▶ Always hand resources back properly