



Garbage Collection

Tobias Wrigstad

Slides (slightly mutilated) from Vitaly Shmatikov

Major Areas of Memory

- ▶ Static area
 - Fixed size, fixed content, allocated at compile time
- ▶ Stacken (one per thread/process)
 - Variable size, variable content (activation records)
 - Used for managing function calls and returns
- ▶ Heap
 - Fixed size, variable content
 - Dynamically allocated objects and data structures
 - Examples: ML reference cells, malloc in C, new in Java

Cells and Liveness

- ▶ Cell = data item in the heap
 - Cells are “pointed to” by pointers held in registers, stack, global/static memory, or in other heap cells
- ▶ Roots: registers, stack locations, global/static variables
- ▶ A cell is live if its address is held in a root or held by another live cell in the heap


Garbage

- ▶ Garbage is a block of heap memory that cannot be accessed by the program
 - An allocated block of heap memory does not have a reference to it (cell is no longer “live”)
 - Another kind of memory error: a reference exists to a block of memory that is no longer allocated
- ▶ Garbage collection (GC) - automatic management of dynamically allocated storage
 - Reclaim unused heap blocks for later use by program


Example of Garbage

```
class Node {  
    int value;  
    Node next;  
}  
Node p, q;
```


```
p = new Node();  
q = new Node();  
q = p;  
delete p;
```



(a)



(b)



(c)

Why Garbage Collection?

- ▶ We have heaps of memory available to our programs, which they use
- ▶ ... badly
 - Memory leaks, dangling references, double free, misaligned addresses, null pointer dereference, heap fragmentation
- ▶ **Also:** Explicit memory management breaks high-level programming abstraction
 - I must expose how a function works internally so that its users can manage memory correctly

GC and Programming Languages

- ▶ GC is not necessarily a language feature
- ▶ "A pragmatic concern for automatic and efficient heap management"
 - GC'd PLs: Lisp, Scheme, Prolog, Smalltalk ...
 - Manual MM: C, C++, Objective-C
 - But garbage collection libraries have been built for them
- ▶ GC revival started in the 90's
 - Object-oriented languages: Java, Python, Ruby, C#
 - Functional languages: ML, Haskell, Erlang, Clojure

The Perfect Garbage Collector

- ▶ No visible impact on program execution
- ▶ Works with any program and its data structures
- ▶ Manages the heap efficiently
 - Always satisfies an allocation request and does not fragment

Summary of GC Techniques

- ▶ Reference counting
 - Directly keeps track of live cells
 - GC takes place whenever heap block is allocated
 - Doesn't detect all garbage
- ▶ Tracing
 - GC takes place and identifies live cells when a request for memory fails
 - Mark-sweep
 - Copy collection
- ▶ Modern techniques: Generational GC

Reference Counting


- ▶ Simply count the number of references to a cell
- ▶ Requires space and time overhead to store the count and increment (decrement) each time a reference is added (removed)
 - Reference counts are maintained in real-time, so no “stop-and-gag” effect
 - Incremental garbage collection
- ▶ Unix file system uses a reference count for files
- ▶ C++ “smart pointer” (e.g., `auto_ptr`) use reference counts

Spot the memory leak(s)

```
-(void) setName: (NSString*) _name {  
    [name release];  
    [_name retain];  
    name = _name;  
}
```

```
/* In the class Node */  
-(void) connect: (Node*) n {  
    if (node == nil) {  
        [n retain];  
        node = n;  
        [node connect: self];  
    }  
}
```

Reference Counting: Example




Reference Counting: Strengths

- ▶ Incremental overhead
 - Cell management interleaved with program execution
 - Good for interactive or real-time computation
- ▶ Relatively easy to implement
- ▶ Can coexist with manual memory management
- ▶ Spatial locality of reference is good
 - Access pattern to virtual memory pages no worse than the program, so no excessive paging
- ▶ Can re-use freed cells immediately
 - If $RC == 0$, put back onto the free list


Reference Counting: Weaknesses

- ▶ Space overhead
 - 1 word for the count
- ▶ Time overhead
 - Updating a pointer to point to a new cell requires:
 - Check to ensure that it is not a self-reference
 - Decrement the count on the old cell, possibly deleting it
 - Update the pointer with the address of the new cell
 - Increment the count on the new cell
- ▶ One missed increment/decrement results in a dangling pointer / memory leak
- ▶ Cyclic data structures may cause leaks
- ▶ Very expensive in concurrent II parallel programs


Reference Counting: Example



Reference Counting: Cycles




Reference Counting: Cycles




Mark-Sweep Garbage Collection

- ▶ Each cell has a mark bit
- ▶ Garbage remains unreachable and undetected until heap is used up; then GC goes to work, while program execution is suspended
- ▶ Marking phase
 - Starting from the roots, set the mark bit on all live cells
- ▶ Sweep phase
 - Return all unmarked cells to the free list
 - Reset the mark bit on all marked cells


Mark-Sweep Example (1)




Mark-Sweep Example (2)



Mark-Sweep Example (3)



Mark-Sweep Example (4)




Mark-Sweep Costs and Benefits

- ▶ **Good**: handles cycles correctly
- ▶ **Good**: no space overhead
 - 1 bit used for marking cells may limit max values that can be stored in a cell (e.g., for integer cells)
- ▶ **Bad**: normal execution must be suspended
- ▶ **Bad**: may touch all virtual memory pages
 - May lead to excessive paging if the working-set size is small and the heap is not all in physical memory
- ▶ **Bad**: heap may fragment
 - Cache misses, page thrashing; more complex allocation

Copying Collector

- ▶ Divide the heap into “from-space” and “to-space”
- ▶ Cells in from-space are traced and live cells are copied (“scavenged”) into to-space
 - To keep data structures linked, must update pointers for roots and cells that point into from-space
 - This is why references in Java and other languages are not pointers, but indirect abstractions for pointers
 - Only garbage is left in from-space
- ▶ When to-space fills up, the roles flip
 - Old to-space becomes from-space, and vice versa

Copying a Linked List [Cheney's algorithm]



Flipping Spaces


to-space



pointer




forwarding address



from-space

root →



Copying Collector Tradeoffs

- ▶ Good: very low cell allocation overhead
 - Out-of-space check requires just an addr comparison
 - Can efficiently allocate variable-sized cells
- ▶ Good: compacting
 - Eliminates fragmentation, good locality of reference
- ▶ Bad: twice the memory footprint
 - Probably Ok for 64-bit architectures (except for paging)
 - When copying, pages of both spaces need to be swapped in. For programs with large memory footprints, this could lead to lots of page faults for very little garbage collected
 - Large physical memory helps

Generational Garbage Collection

- ▶ Observation: most cells that die, die young
 - Nested scopes are entered and exited more frequently, so temporary objects in a nested scope are born and die close together in time
 - Inner expressions in Scheme are younger than outer expressions, so they become garbage sooner
- ▶ Divide the heap into generations, and GC the younger cells more frequently
 - Don't have to trace all cells during a GC cycle
 - Periodically reap the “older generations”
 - Amortize the cost across generations

Generational Observations

- ▶ Can measure “youth” by time or by growth rate
- ▶ Common Lisp: 50-90% of objects die before they are 10KB old
- ▶ Glasgow Haskell: 75-95% die within 10KB
 - No more than 5% survive beyond 1MB
- ▶ Standard ML of NJ reclaims over 98% of objects of any given generation during a collection
- ▶ C: one study showed that over 1/2 of the heap was garbage within 10KB and less than 10% lived for longer than 32KB

Frågor?





Genomgång projektarbete

Kommentarer efter TF:s föreläsning

- ▶ Gruppsammansättning
- ▶ Dagliga möten
- ▶ Låt inte saker ruttna
- ▶ Kontinuerlig integration
- ▶ Att alltid ha ett levande system
- ▶ T.ex. att själva använda ert eget system för att hitta buggar

Uppgiften

- Att skriva en egen minneshanterare i C

Manuell minneshantering med explicit allokering och avallokering

Referensräkning med explicita modifieringar av referensräknaren

Tracing GC i med en enkel mark-sweep-algoritm



16-sidig specifikation i repot

specifikation

”kurskrav”

Manuell minneshantering

- alloc
- free
- justering av allocs beteende
 - snabb allokering som fragmenterar
 - mindre fragmentering
 - hög grad av lokalitet

Referensräkning

- alloc (refcount = 1)
- retain
- release
- count

Mark-sweep

- alloc
- automatisk gc vid slut på minne
- Hitta rotpekare på stacken (kod finnes)
- Scanna minnet efter möjliga pekare

Allokering med "metadata"

- `alloc("***d") = alloc(3 * sizeof(void*) + sizeof(double))`

* = pekare

c = char

d = double

f = float

i = integer

l = long integer

- Metadatat kan användas för att slippa scanna heapen

Hur arbetet skall utföras

- Ni är indelade i team om 6 personer

Det kommer att vara viss flux denna vecka (alla är inte aktiva, etc.)

Teamen beskrivs på kurswebben under "grupper"

- Varje team delas in 3 programmeringspar

Dessa skall rotera varje vecka så att man får jobba med (nästan) alla

I vissa grupper kommer det att blir ett "par" med tre programmerare

- Utgå från Scrum för skapandet av er process

<http://sv.wikipedia.org/wiki/Scrum>

- **DEADLINE: 17/12**

Projektet

- All kod skall skrivas i C

Testkod etc. får gärna skrivas i annat språk

- Testdriven utveckling skall användas (se bilder från Justin)
- Obligatoriska verktyg

Trello (för planering och uppföljning av arbete)

GitHub (för versionshantering och issue tracking)


git

gcc

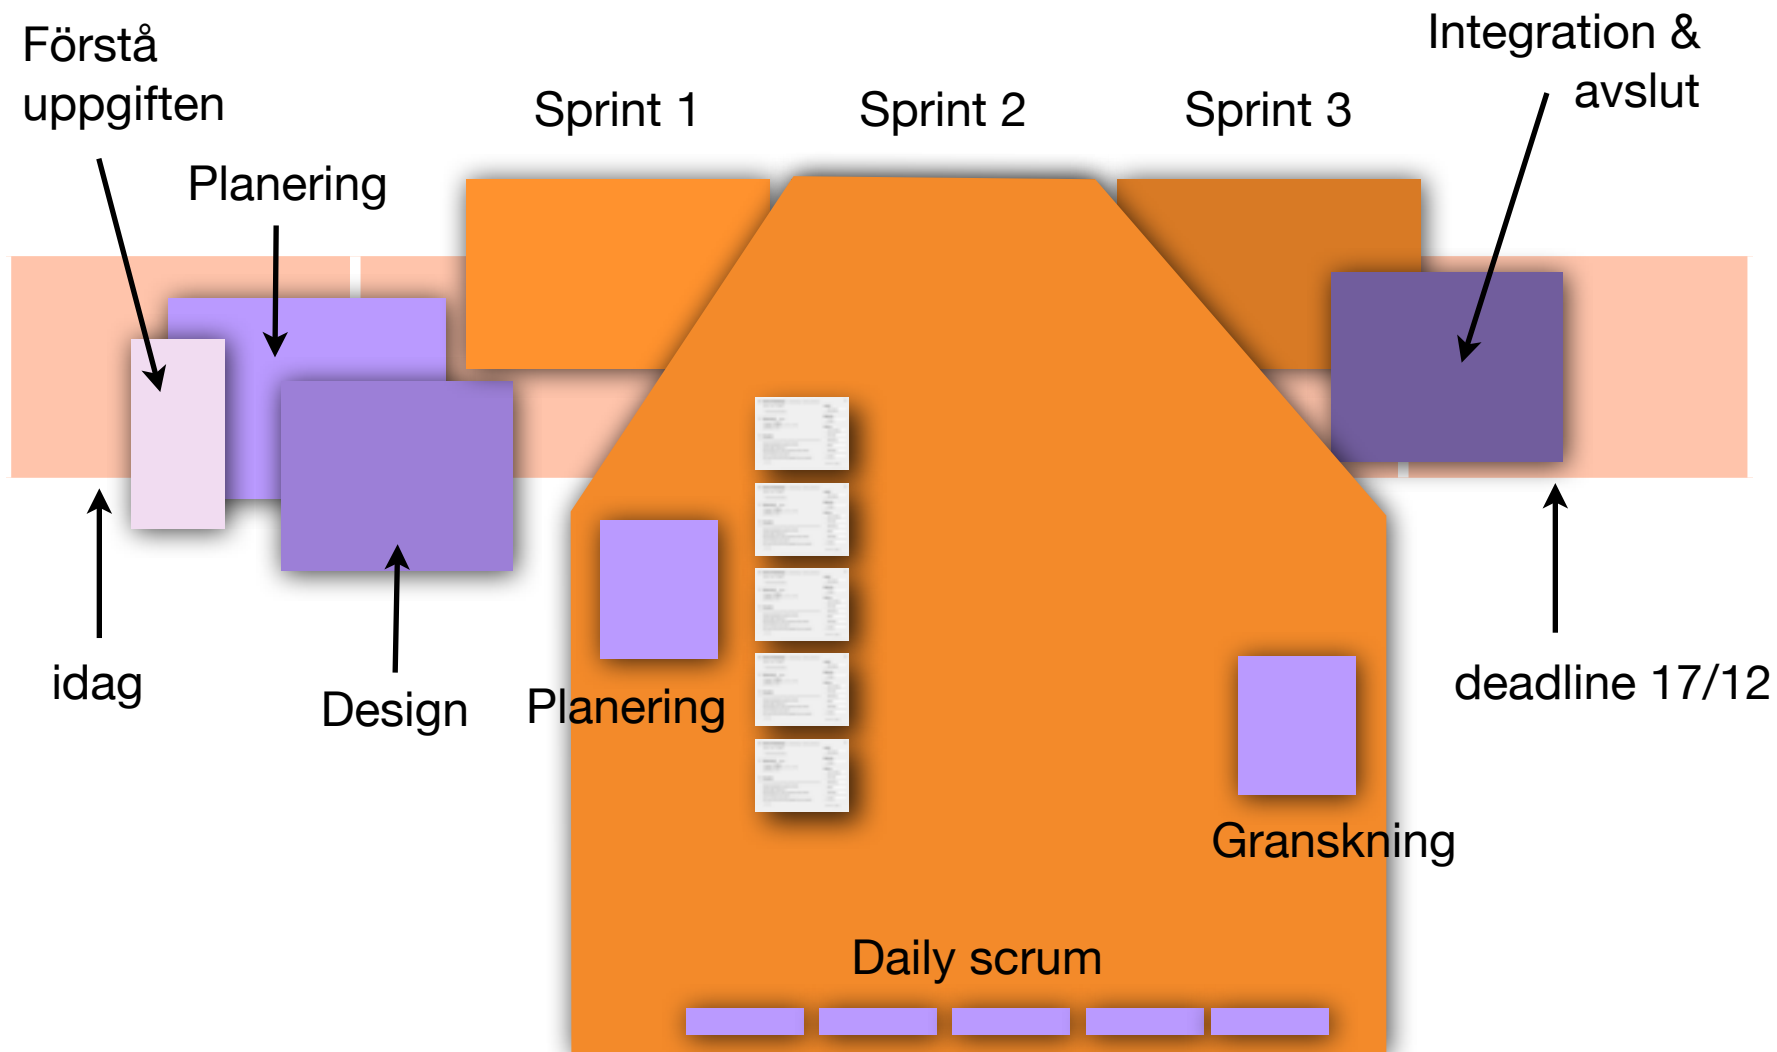
Make

valgrind

Tidplan (exempel)



Tidplan (exempel)



Vad skall lämnas in?

- Projektdagbok

Vem har jobbat på vad, arbetstider per deluppgift, kort reflektion kring projektet, motivation av designen

- Övergripande designdokument

Vilken uppdelning i delsystem, hur sitter delsystemen ihop?

- Koddokumentation på gränssnittsnivå

Á la Javas API

- Gränssnitten mellan modulerna

Det centrala dokumentet för parallell utveckling!

- Själva koden

- Tester

Plattformsoberoende

- Programmet som ni utvecklar skall vara plattformsoberoende

Specifikt skall det fungera både på X86 och SPARC på IT

- Vissa aspekter av minneshanteringen är plattformsspecifika!

T.ex. åt vilket håll stacken växer

Pointer alignment?

Spelar big/little endian någon roll?

Bedömning

- Den slutinlämnade kodens kvalitet och kompletthet
- Inlämnad dokumentation
- Kvalitet på egna testfall
- Aktivt deltagande i utvecklingsprocessen
- Projektdagboken
- **Beräknad arbetsåtgång: ca 133 arbetstimmar per person**

Vad händer om man inte är klar?

- Syftet med uppgiften är inte att producera ett färdigt system
- ...*därmed inte sagt att det inte är viktigt att producera ett färdigt system*
- Om man inte är klar 17/12 kan tre saker hända:

Man blir ändå **godkänd** eftersom man har dokumenterat sina brister, har skrivit bra kod i övrigt, och har en trovärdig plan för resten av systemet

Man får **rest** och ett nytt sista inlämningdatum allra senast 10/1 –14

Man har inte gjort ett seriöst försök och får därför **underkänt** och är välkommen igen HT2014

Får man redovisa mål som vanligt under projektet?

- JA! Det är en av poängerna!
- Vi byter nu till 1 labb/vecka (onsdagar 13–17)
- Syftet med dessa labbar är att man skall kunna redovisa mål som vanligt, samt få hjälp med projektet
- Notera att man får redovisa mål X65 även på labbarna

Konton

- trello.com — skapa ett eget konto
- github.com — skapa ett eget konto
- För att få tillgång till gruppens privata repo, maila Elias (elias.castegren@it.uu.se)

Maila endast en gång per grupp!

Ni skall använda de privata repo som vi ger ut — historiken där är viktig!